

# Adaptive Concurrency Control

*Despite the Looking Glass,  
One Concurrency Control Does Not Fit All*

---

AARON J. ELMORE



DIXIN "TOTEM" TANG



HAO JIANG























# Choosing the right specialist for a task is important.

---

WE NEED MORE THAN JUST SPECIALIZATION





















# Building Many Great “Specialists”

---

Specialized work in the last decade figuring out how to make OLTP DBs go fast.

H-Store, Hekaton,  
Silo, Doppel,  
VLL, Orthus, TicToc,  
Foedus, MOCC,  
[Insert your favorite]



Despite the looking glass,  
no one concurrency control protocol can fit all.





# Excel Under the Right Circumstances

---

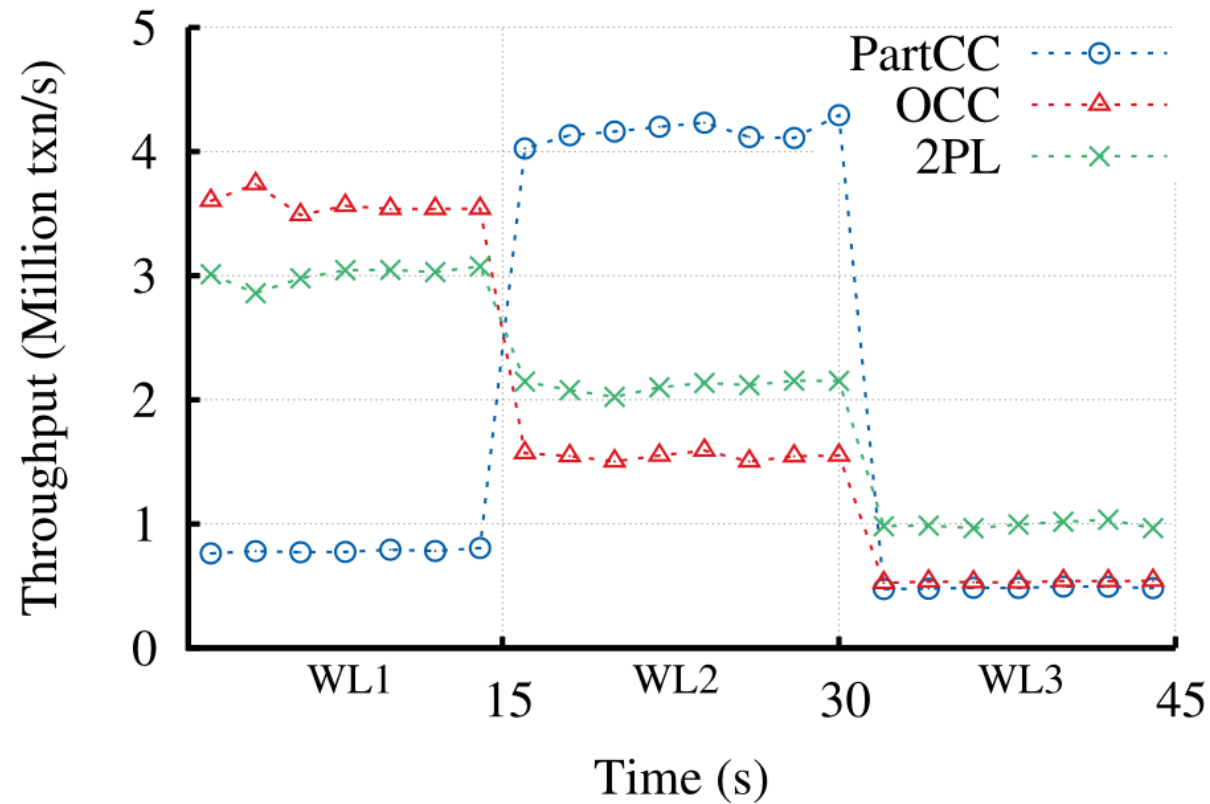
H-Store's single threaded (**PartCC**) model suffers with non-partitionable workloads.

Optimistic CC (**OCC**) suffers with high write contention.

Pessimistic CC (**2PL**) is slower than others with low skew.



# Sample Workload Variation



WL1: Not partitionable, low skew

WL2: Partitionable

WL3: Not partitionable, high skew

YCSB on main-mem DB

OCC  $\approx$  Silo

2PL  $\approx$  VLL & No-Wait

PartCC  $\approx$  Single Threaded

**Each CC excels in different scenarios.**

Many have observed this effect to differing degrees, e.g. Carey et al.



Can we take advantage of specialists (*a CC protocol*) without being brittle to less-than-ideal workload?

---



Yes we can!

... but which specialists and how many?

---




# Even with Specialists

---

Specializing is good, but people (DBAs) will struggle to pick the ideal CC.

- Workloads are not known a priori
- Sudden shifts in skew happen
- Gradual changes can occur



# Our Vision

---

Many great static solutions/specialists.

Every ~5 years more self-managed, autonomous, etc. DBMS. We need more than some form of control theory, auto-tuning, or parameter setting.

DBMS systems must be able to adaptively  
**run with multiple modes/specialists concurrently.**



# Step towards this Vision: Adaptive Concurrency Control

---

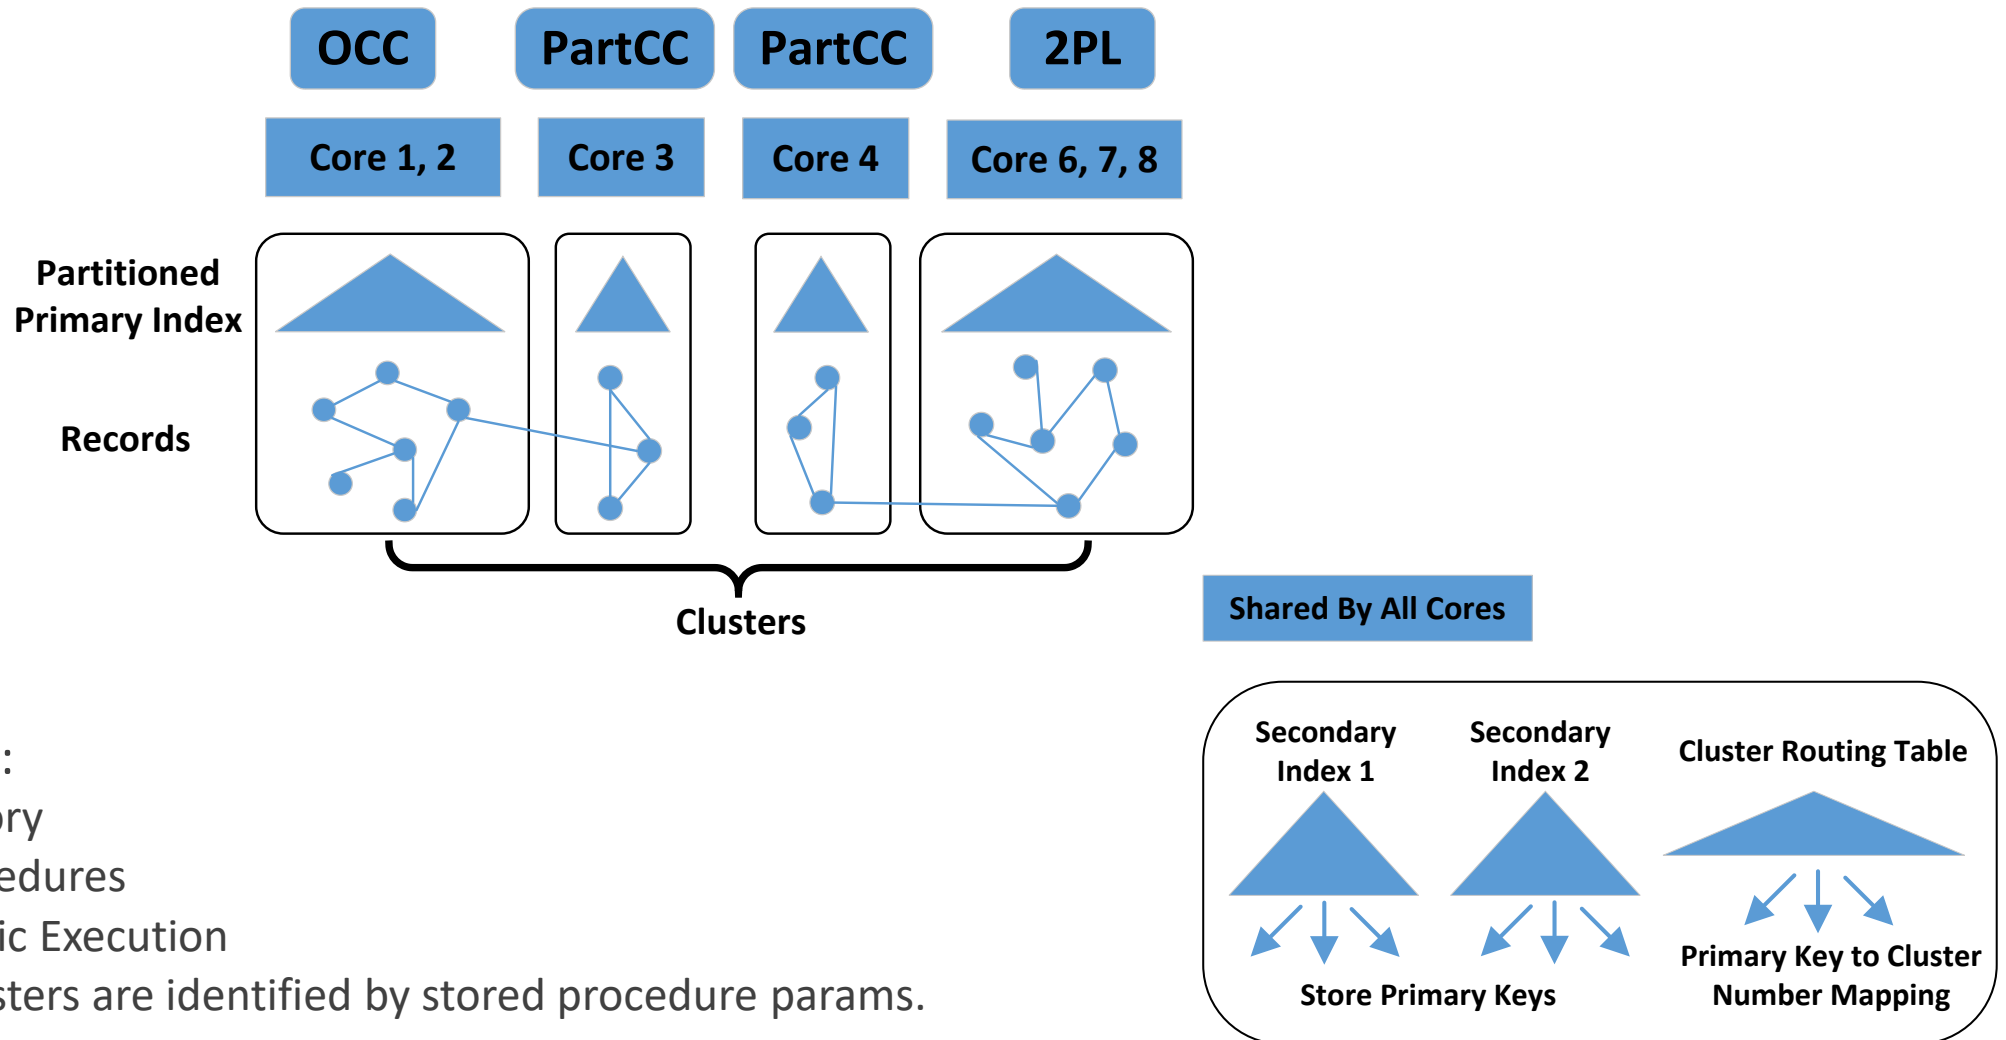
An adaptive and high performance transactional database that:

**Clusters data automatically (partitioning with core assignment)**

**Adaptively selects a concurrency control per cluster**

**Supports a general framework for mixing CC protocols**

# ACC Overview



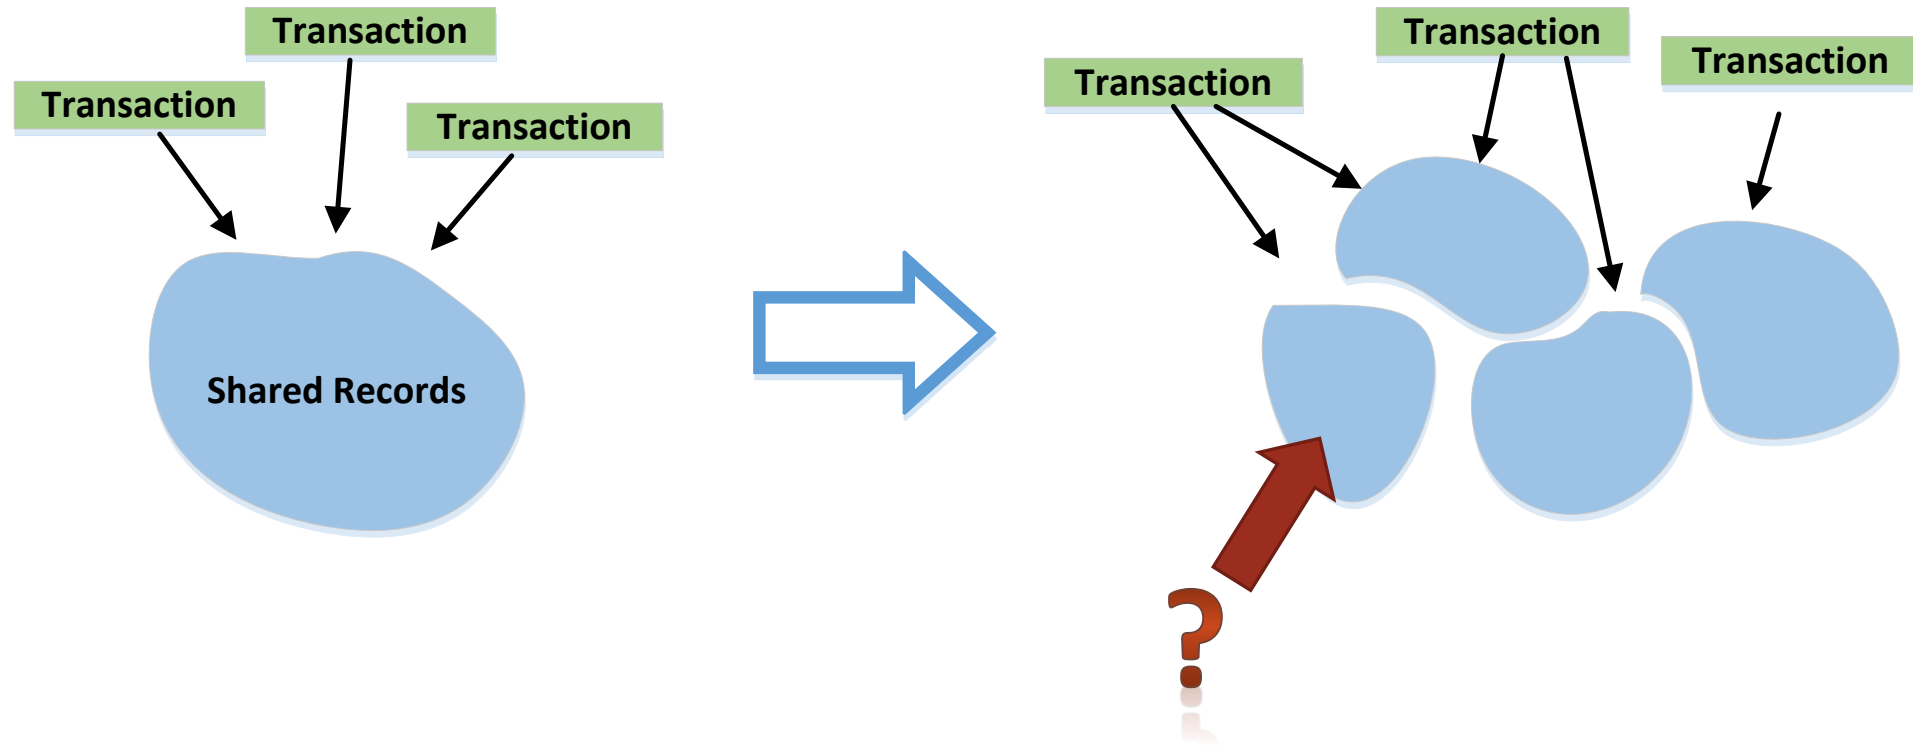


# Adaptive CC Selection

---

# Choosing a CC Protocol

---



For now assume we have a good technique to map data into clusters and assign cores / workers.



# CC Selection is More than BB ML

---

**Goal:** For a given cluster, pick the CC with the highest throughput for a workload

**Approach:** Learn a model to select right CC (this isn't the 80s)

## **Challenges:**

- Feature engineering
- Lightweight feature extraction
- Training models


# Workload Modeling

---

Identify features that have the biggest impact on the comparative performance of various CCs.

Look for signals that impact:

- Partitionability
- Record Skew
- Write Contention



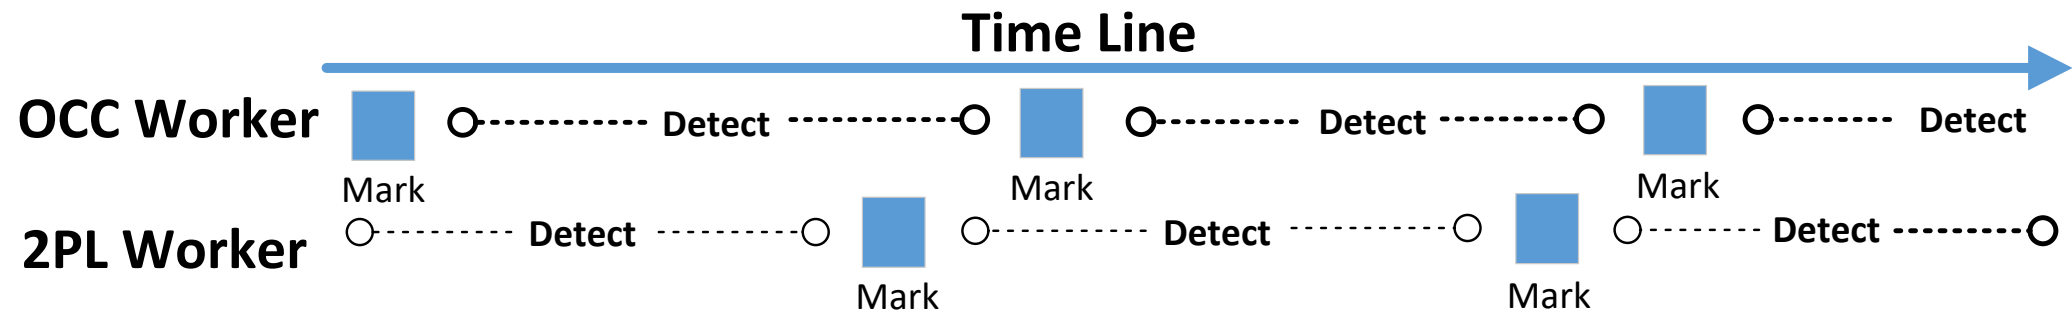


# Lightweight Feature Extraction

Feature extraction at scale and across cores is difficult.

- Different CCs exhibit different behaviors
- Centralized solution can be limiting

Think about detecting record conflicts / skew across CC

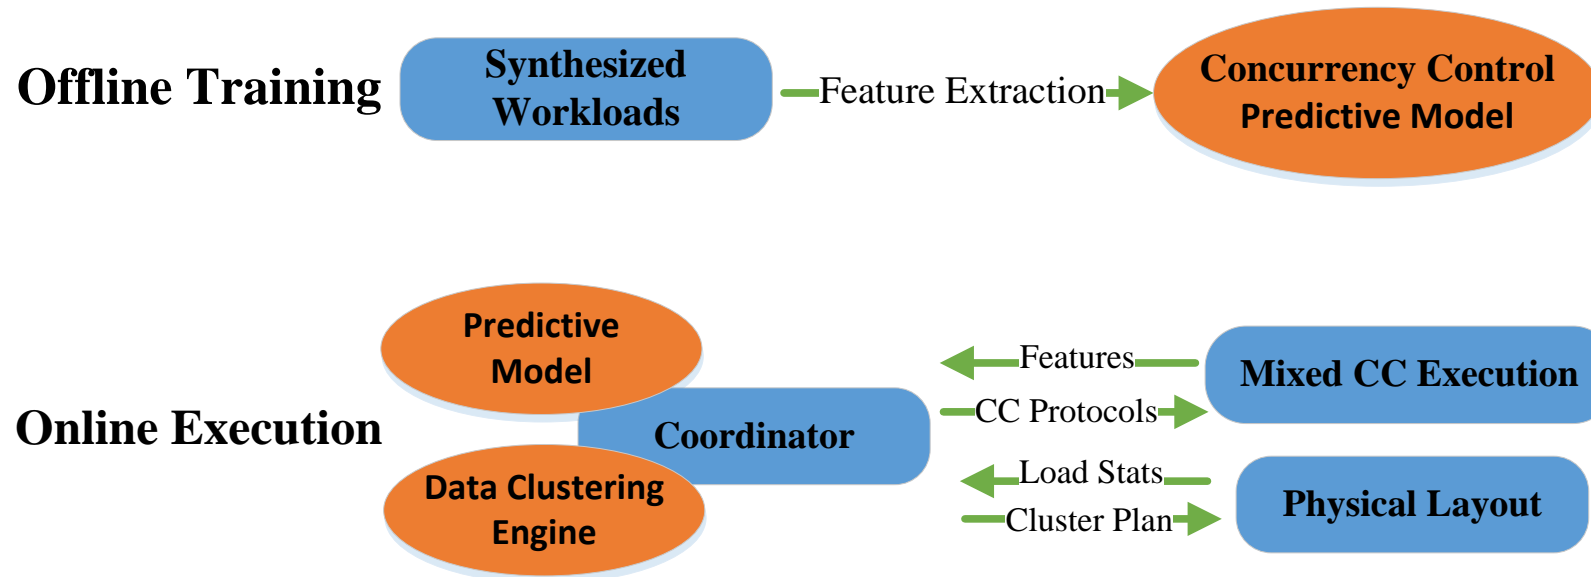


# Putting CC Selection All Together

---

Training the model

Using the model online





# Mixing CC Protocols

---

# Mixing CC is More Than Using Locks

---

**Goal:** Support transactions that span multiple clusters with different CC protocols with correctness *and* performance.

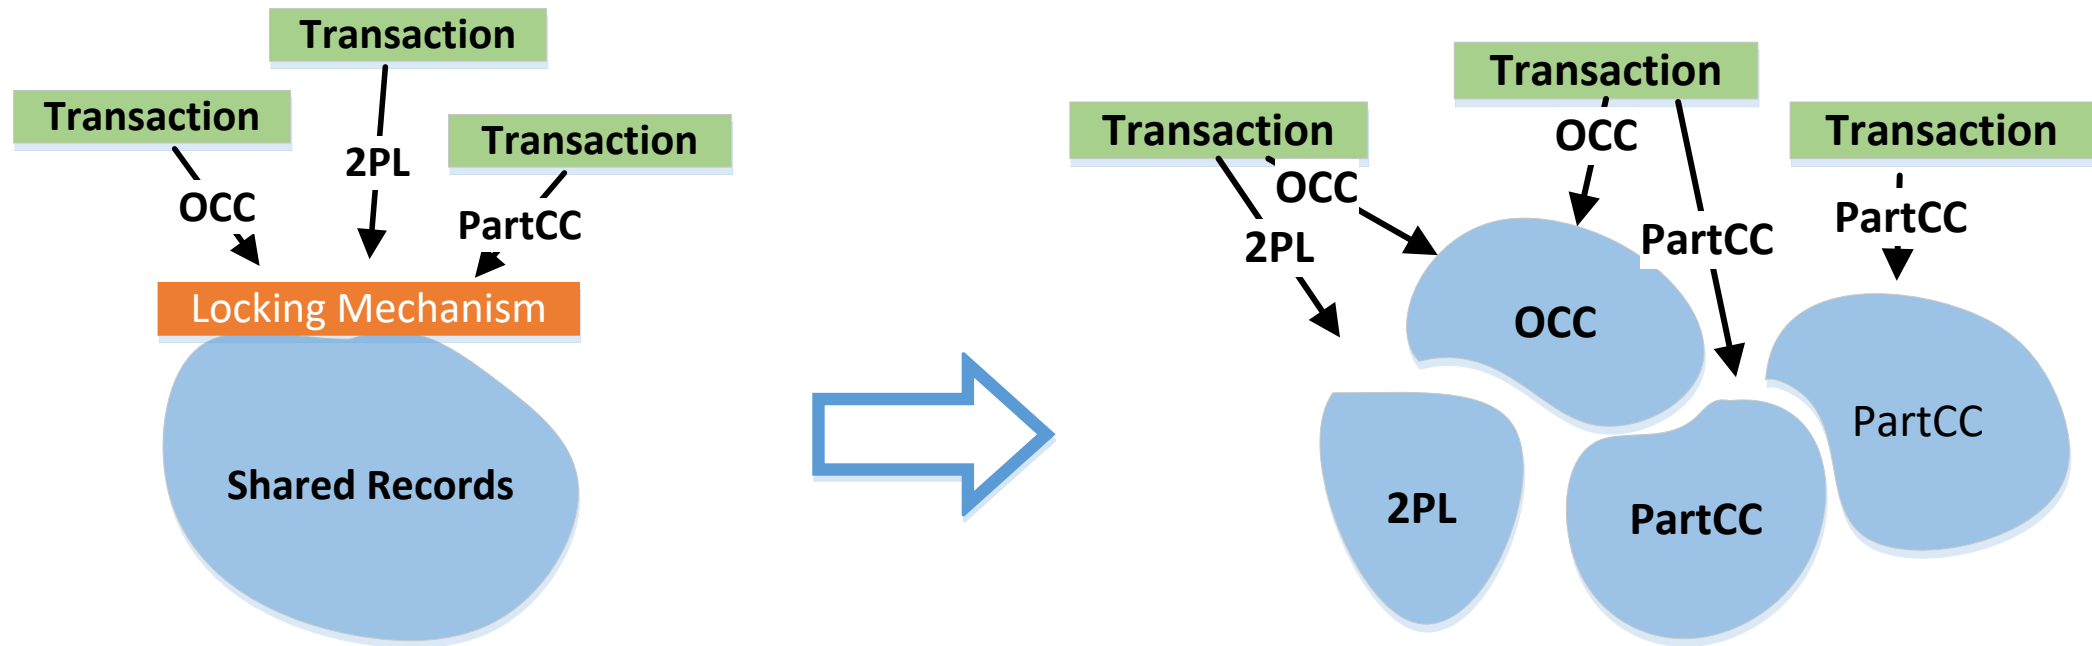
**Approach:** Use the record's “native” CC and span multiple CCs.

## Challenges:

- No modification of existing CC protocol
- Lightweight synchronization for conflict detection across protocols
  - (existing approaches rely on centralized or global locking)
- Maintain transaction properties of underlying protocols (e.g. conflict **serializability**, deadlock-free, strictness, etc).



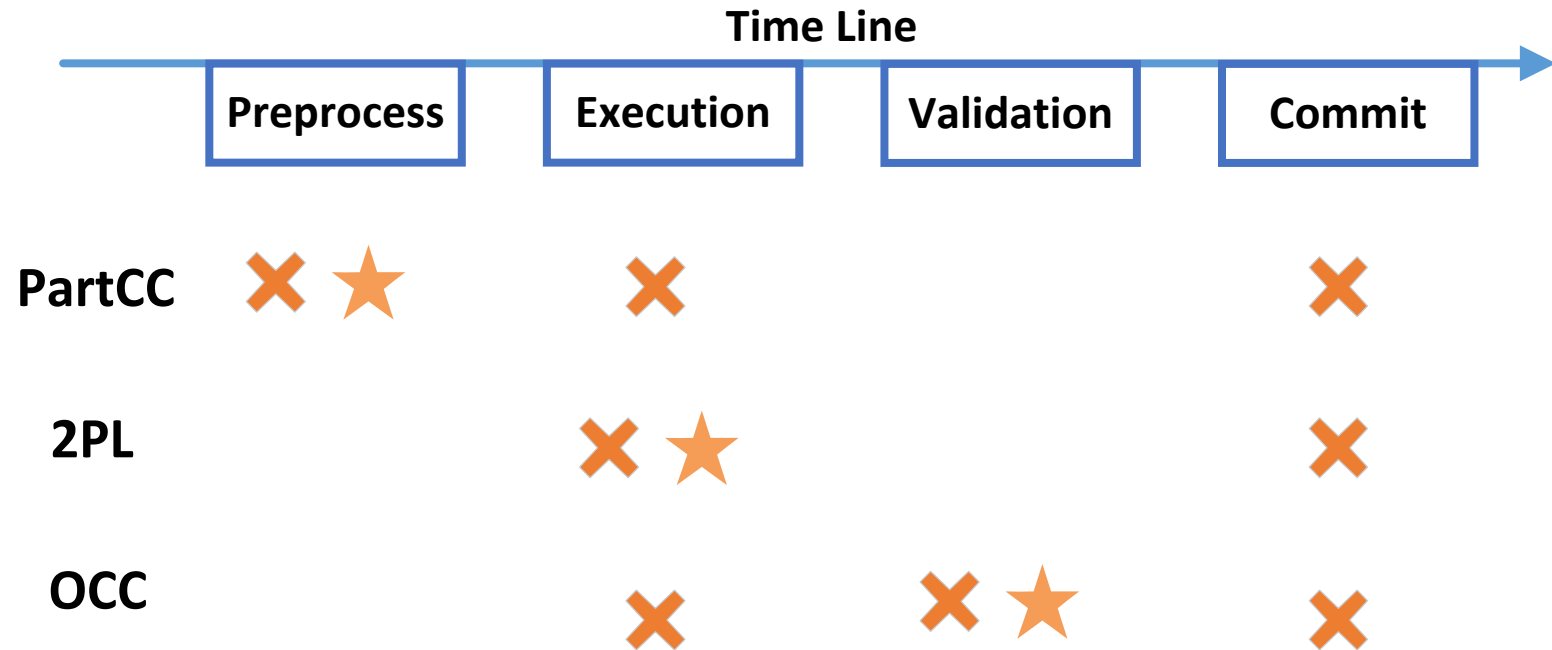
# Data-oriented mixed Concurrency Control (DomCC)



Instead of setting global CC

Use the CC “belonging to the tuple”

# DomCC for PartCC, 2PL, and OCC



 The protocol execution includes this phase

 The protocol makes transactions wait due to conflicts in this phase

# Online CC Reconfiguration

---



# Online CC Reconfiguration

---

**Goal:** Allow clusters to change CC without stopping execution.


**Approach:** Allow for hybrid CC during change window.

## **Challenges:**

- Lightweight synchronization for view of CC amongst workers
- General framework for switching
- Partitioned & Non-Partitioned Structures

# Online Protocol Switching

Challenge to switch protocols without stopping the entire system or heavyweight synchronization.



# Data Clustering

---



# Clustering is More than Partitioning

---

**Goal:** Maximize throughput and load balance as secondary

**Approach:** **Holistic approach in progress**

**Challenges:**

- Unknown/flexible number of partitions
- Core allocation (flexible number of workers/capacity)
- Synchronization overhead of CC
- Cost of inter and intra CC distributed TXNs

# Lots of Related Work We Skipped Today

---

Studying CC performance

Hybrid CCs

Optimizing CC

Auto-tuning, self-managed, and adaptive DBs

Hybrid Systems

# Rich Set of Future Problems with ACC

---

Clustering Algorithms

Adaptive Replication

Online Training of Models

Mixed Logging Protocols

Partitioning of Indices (primary and secondary)

Shrinking to Min Capacity



Let's rethink DBMS architecture  
to support a mix-n-match  
approach for components.

---

...AND FIGURE OUT HOW TO MAKE THEM PLAY NICELY TOGETHER